



# Enabling Reproducibility with Docker

# Enabling Reproducibility with Docker

---

- Gerardo Hidalgo-Cuellar
- *RC Homepage:* <https://www.colorado.edu/rc/>
- *RC Docs:* <https://curc.readthedocs.io/en/latest/>
- *RC Helpdesk:* [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Course Materials:  
[https://github.com/ResearchComputing/Containers\\_Spring\\_2022](https://github.com/ResearchComputing/Containers_Spring_2022)

# Outline

---

- Part 1: Container fundamentals and Docker (4/7/2022)
  - Reproducibility and the Case for Containers
  - Containers
  - Docker
    - Images and Containers
    - Commands
    - File Access
    - Building Docker Images
    - DockerHub
- Part 2: Containers for HPC w/ Singularity (4/14/2022)

# Tutorial Files:

---

- This tutorial will have interactive components
- If you would like to participate in the demos provided for this tutorial then first clone the test files from GitHub to your desired location:
  1. Navigate to a desired directory
  2. Clone the repository:

```
git clone https://github.com/ResearchComputing/Containers\_Spring\_2022.git
```
  3. Navigate into the directory and store the path into a variable:

```
cd Containers_Spring_2022  
export CONTAINER_ROOT=$(pwd)
```



# Reproducibility and Research

---

- Scientific Software is often challenging to work with
  - Difficult installation
  - Low support from the developers
  - Very outdated
  - Complex Dependency trees
- Because of this it's often desired for a software to be repeatable and accurate.
- *But installs are only done once. Why should I care about reproducible applications.*

# Reproducibility and Research

---

- Research is Collaborative
  - Team members work together to get projects done.
  - Reproducibility ensures all members of a team can provide productivity towards a project.
- Research is Correcting
  - Research is hard
  - Academic reviews are commonplace
  - Someone may wish to accurately reproduce your work
- Research is Continuous
  - You may be working on a single project for a long period of time
  - What happens in you move, but bring your work to another system?

# Options for reproducibility

---

- Lots of options!
  - Detailed instructions
  - Software bundles
  - Virtual Environments
    - Python, Anaconda, Spack
- But do they really enable accurate reproducibility?
  - Incorrect installs?
  - Hardware or OS?
  - Performance?

# Containers

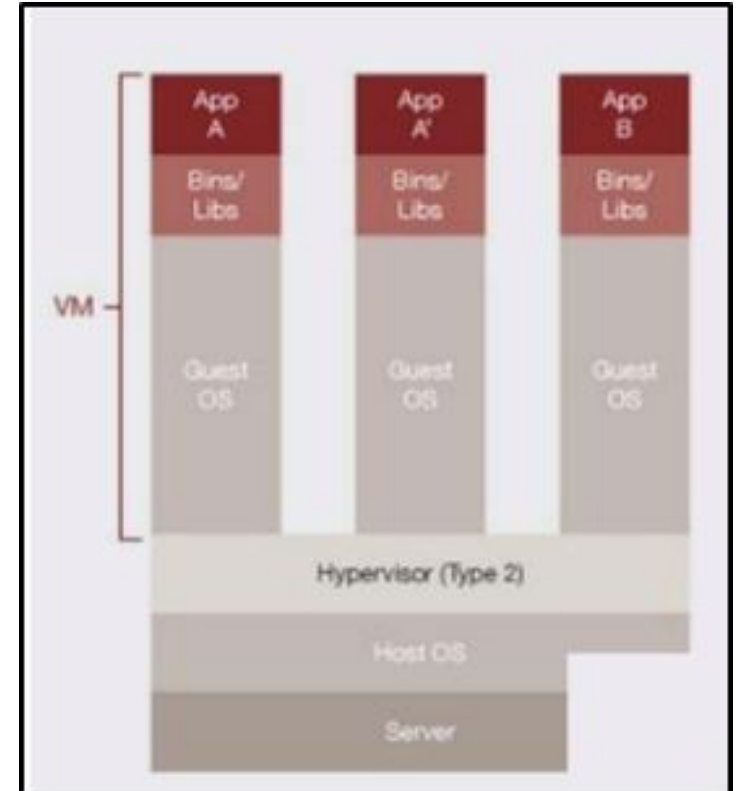
---

- A Container is an isolated environment: packaged bundle of libraries, dependencies, and files that runs as a process under a host OS
- Containers use an application on the host operating system called a Container Manager
  - Manages operating system and libraries run as containers
  - Similar to virtual machines, but does not need dedicated CPUs memory or storage



# Virtualization (1)

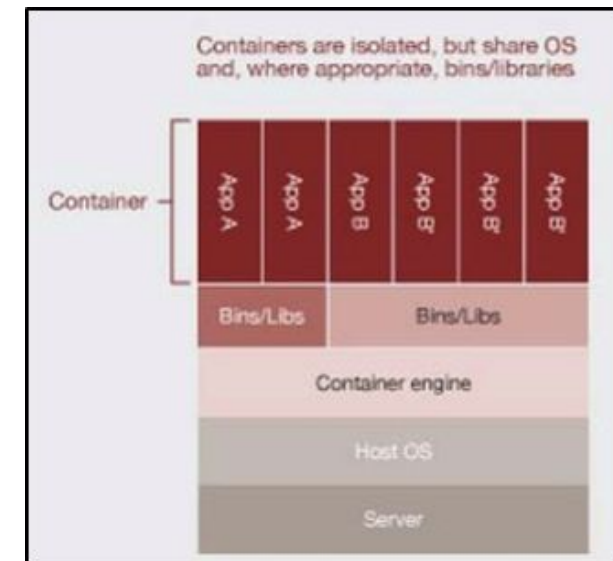
- Virtualization is a technology that utilizes software to abstract components of a technology
- Common application is Hardware Virtualization
  - Virtual Machines
    - Partitions off Memory, CPU, GPU, and Storage
    - Runs a virtual OS
    - Runs software on the virtualized machine
  - *Examples: VMware, Virtualbox*



Material courtesy: M. Cuma, U. Utah

# Virtualization (2)

- Another use of virtualization is in OS Level Virtualization
  - Can run many isolated guest OS instances under a host OS kernel
  - This virtualization is what is used by Docker and other container software.
  - Virtualizing *software*, not *hardware*
  - Share a kernel
  - Best of both worlds!
    - Isolated environments
    - No hardware partitioning



Material courtesy: M. Cuma, U. Utah

# Containerization Software

---

- **Docker**
  - Well established – largest user base
  - Has Docker Hub for container sharing
  - Problematic with HPC (Fix incoming!)
- **Singularity**
  - Designed for HPC
  - Second largest user base
  - Developed for scientific use
- **Charliecloud; Shifter**
  - Designed for HPC
  - Based on Docker
  - Less user-friendly



# Installing Docker

---

- Docker Desktop
  - Comfy GUI to help keep track of containers and images!
  - Available on all operating systems (*beta on Linux*)
  - Windows users can enable WSL2 support following the instructions here:  
<https://docs.docker.com/docker-for-windows/install/>
- Docker hosted lab environment (Need Docker account, limited availability)
  - <https://labs.play-with-docker.com/>
- Docker toolbox
  - Legacy solution for Windows and Mac for versions that do not meet the version requirements.
  - Utilizes the Virtual Box hypervisor for virtualization

# Docker: 3 main components

---

- **Docker File**

- Like DNA, code that tells docker how to build an image

- **Image**

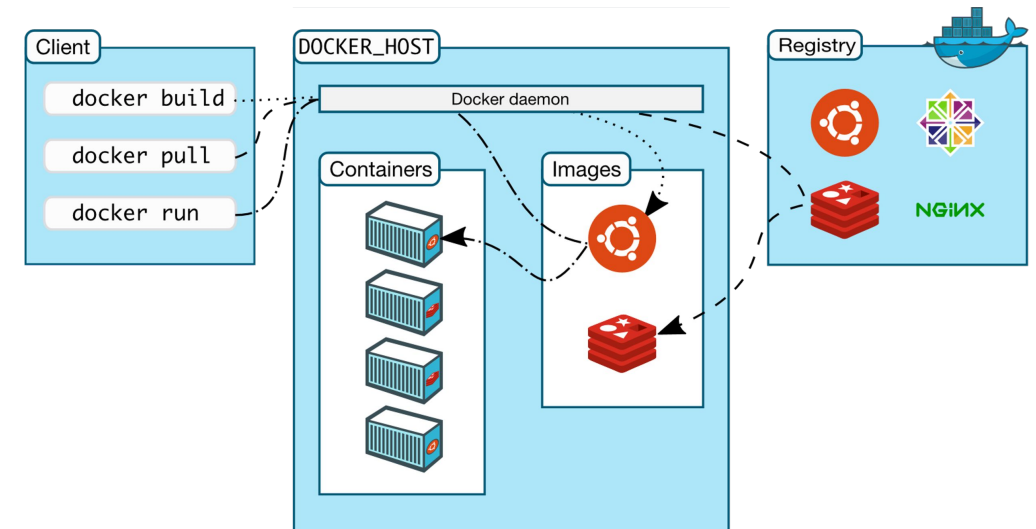
- Snapshot of your software along with all of its dependencies (down to OS level)
- Immutable (mostly) and can be used to spin up multiple containers

- **Container**

- Running instances of images that are isolated and have their own sets of environments and processes
- Actual software running in the real world

# Docker Nuts and Bolts

- Docker runs on a concept of images and containers.
  - **Images:** Saved snapshots of a container environment.
    - Made from a Dockerfile or pulled from Docker Hub
    - Stored in the Docker cache on your disk
    - Immutable (mostly...)
  - **Containers:** Instances of images that are generated by Docker when an image is 'run'
    - Instance of image running in memory
    - Ephemeral and state cannot be saved
    - Can be run interactively





# Docker 'Hello World'

---

- Let's start with something simple:

- Docker "Hello, World!"
- Relatively small image
- No dependencies
- Built as a general test case

- Command we will run:

```
docker run hello-world
```

# Docker Commands

---

- Docker Commands are usually in the form of:  
`docker <sub-command> <flags> <target/command>`
- Examples:  
`docker run -it myimage`  
`docker container ls`  
`docker image prune`

# Launching a Docker Container

---

- Launch docker image as a container:

```
docker run <image-name>
```

- Run a docker image interactively:

```
docker run -it <image-name>
```

- If an image is not on the system, then Docker will search DockerHub to see if the image exists, and pull it down locally
- Specify commands after your image to execute specific software in your container.

```
docker run <image-name> <program>
```

Example:

```
docker run -it ubuntu bash
```

# Listing Containers

---

- When a container is run it is assigned a name, an ID, name of the image used to run the container, current status.

- List all currently running containers

`docker ps`

- List all containers

`docker ps -a`

# Stopping/Removing a container

---

- You can stop a container using the “stop” command:

```
docker stop <name or ID>
```

If you don't know the docker name or id you can list containers

- If you don't want a stopped container taking up space you can remove it with the remove command:

```
docker rm <name or ID>
```

```
docker container rm <name or ID>
```

# Images

---

- To see a list of images (templates) for our container:

```
docker images  
docker image ls
```

- To remove an image:

```
docker rmi <name>  
docker image rm <name>
```

- To download an image but *not* run an image:

```
docker pull <name>
```



# Exploring a Docker Container

---

- Docker containers are running tiny operating systems!
- We can explore the operating system by invoking a shell  
`docker run -it ubuntu bash`
- This command launches the ubuntu Docker container with the command 'bash'
- We can also run a command on an already running container with the "docker exec" command:  
`docker run -d ubuntu sleep 100`  
`docker exec <container-id> cat /etc/*release*`

---

# Demo 1: Running a Container

# Demo 1: GROMACS

---

- GROMACS is a molecular dynamics application that can often be a complex and challenging installation for the average user. Linux and Mac only
  - Dense Documentation
  - Software requires compilation
- Luckily, this can be trivialized with Docker!
  - Run the command:

```
docker run gromacs/gromacs gmx help commands
```

```
docker run -it gromacs/gromacs
```

# Demo 1: GROMACS

---

- An example using *pdb2gmx* from the tutorial [KALP15 in DPPC](#):

```
$ mkdir $HOME/data ; cd $HOME/data
```

```
$ wget http://www.mdtutorials.com/gmx/membrane\_protein/Files/KALP-15\_princ.pdb
```

```
$ docker run -v $HOME/data:/data -w /data -it gromacs/gromacs gmx pdb2gmx -f KALP-15_princ.pdb -o KALP-15_processed.gro -ignh -ter -water spc
```

- When prompted, choose the GROMOS96 53A6 parameter set (13) and choose "None" for the termini

# Docker Image/Container Commands

---

Container Commands	
<code>docker container ls</code>	List docker containers currently running:
<code>docker container rm &lt;container&gt;</code> <code>docker rm &lt;container&gt;</code>	Remove (an) container(s):
<code>docker container prune</code>	Remove all stopped containers

Image Commands	
<code>docker image ls</code>	List docker images stored in cache:
<code>docker image rm &lt;image&gt;</code> <code>docker rmi &lt;image&gt;</code>	Remove (an) image(s):
<code>docker image prune</code>	Remove unused images

# Docker Image/Container Commands

---

Commands	
<code>docker info</code>	Shows Docker system-wide information
<code>docker inspect &lt;docker-object&gt;</code>	Shows low-level information about an object
<code>docker config &lt;sub-command&gt;</code>	Manage docker configurations
<code>docker stats &lt;container&gt;</code>	Shows container resource usage
<code>docker top &lt;container&gt;</code>	Shows running processes of a container
<code>docker version</code>	Shows docker version information

- More details and commands can be found [on the docker documentation page](#)



# DockerHub

---

- The place where containers live! I.e. Image Registry
- Dockerhub is a Docker hosted library of public and private Docker images, with an account:
  - Free and unlimited public images
  - 1 free private repository
- Great for hosting images for fellow researchers
- “Git-like” commands

# Building a Docker Container

---

- To build a docker container, we need a set of instructions Docker can use to set up the environment.
  - Dockerfile (<- must have this name, no extensions)
- Once we set up our Dockerfile we can use the command  
`docker build -t <image-name> .`
- Then we can run the image with our docker run command  
`docker run <image-name>`

# What's in a Dockerfile

---

- A Dockerfile is simply a text file that contains instructions to build and setup a default Image
  - Commands to build
  - Setting commands
- Requires a source Image
  - a “template” image

```
mtrahan41@MTrahanRazor15:[ ubuntu-gcc ]$ cat Dockerfile
FROM ubuntu:18.04

RUN apt-get update; \
    apt-get install nano -y; \
    apt-get install gcc -y; \
    mkdir target;

WORKDIR /target
```

# What's in a Dockerfile

---

- **FROM**

- start **FROM** a “template” image
- base image gets pulled down from cloud

- **RUN**

- to **RUN** terminal commands
- install dependencies

- **WORKDIR**, **ENV**, etc...

- **CMD**

- execute a default **CMD**

```
mtrahan41@MTrahanRazor15:[ ubuntu-gcc ]$ cat Dockerfile
FROM ubuntu:18.04

RUN apt-get update; \
    apt-get install nano -y; \
    apt-get install gcc -y; \
    mkdir target;

WORKDIR /target
```

---

# Demo 2: Building a Docker Image

# Demo 2: Ubuntu w/ GCC

---

- For this first example we will build a custom Ubuntu image that will provide a location to run the GNU Compiler Collection.
- Dockerfile provided:
- Need to build:
  1. Navigate to the directory:  
`cd $CONTAINER_ROOT/dockerdemo/ubuntu-gcc`
  2. Build the image with:  
`docker build -t test-gcc .`
- Run image as container:  
`docker run -it test-gcc`



# Demo 2: Ubuntu w/ GCC

---

- What happens if we create a file in the container?
- Does it persist if we exit?
  - No! Containers are ephemeral and run in host memory
- How can we persist data?

# Mounting and File Access (1)

---

- So now that we have a working container, how can we access the test files we downloaded?
  - Mounting directories: **Bind Mount**
    - Allows the docker container to access files on the host OS
    - Choose host's *source directory*, files in the directory will be moved to the container's *target directory*
      - **Source Directory**: Directory on the host system. Never within a container.
      - **Target Directory**: Directory in the Docker Container. Never on the host system.
    - A flag set within the docker run command:  
`docker run -v <source-dir>:<target-dir> <image>`

# Mounting and File Access (2)

---

- Mounting directories: Volume Mount
  - Same concept, but volumes are stored within docker cache.
  - Create Docker volumes in your terminal and link your volume directory
  - Similarly linked through the docker run command.

```
docker run -v <volume-name>:<target-dir> <image>
```

# Demo 2 (Cont.): Mounting

---

- Returning to our demo, can we give our container access to our test files?
- Let's use a bind mount!
- In the directory where our Dockerfile lives, use this command (all on one line):

```
docker run -it -v $(pwd)/source:/target test-gcc
```

- Command:

```
gcc hello.c -o hello.exe  
./hello.exe
```

# Modifying a Docker Image

---

- Suppose you have an existing docker image and want to make changes...
  - Rebuild Dockerfile!
  - Usually a bit cumbersome
- No Dockerfile? Use docker commit!
  - First you can run an image interactively and install what you need:

```
$ docker run -it <image-name> bash # or any shell...  
$ apt-get update  
$ apt-get install vim
```
  - Exit, then commit it to a new image

```
$ docker commit <container-id> <new-image-name>
```

# Dockerhub Commands

---

- Download and upload docker images with ease.

```
docker run <image>
```

```
docker pull <image>
```

- Uploading a little more complicated...

- Sign in with:

```
docker login
```

- List docker images with:

```
docker image ls
```

- Tag your image:

```
docker tag <image-id> <your-username>/<image-name>:<tag>
```

- Push!

```
docker push <your-username>/<image-name>
```

---

# Demo 3: NCL container

# Demo 3: NCL Container

---

- For this next example we will be building a Docker image that will run the NCAR Command Language (NCL)
- Dockerfile provided
- Same process:
  1. Navigate to the Dockerfile found at:  
    \$ CONTAINER\_ROOT/dockerdemo/ncl
  2. Build the Dockerfile and name the image: "ncl-demo"
  3. Run "ncl-demo"
- Can we test a sample script?

```
$ docker run -v $(pwd)/source:/target ncl-demo ncl test.ncl
```



# Docker Compose

---

- External Utility that can create and install docker images.
- Builds docker images based on a docker-compose.yml file.
- YAML: YAML Ain't Markup Language
  - Data serialization language
- Describes containers you wish to build with what features.
- Not a docker command but comes bundled with Docker Desktop!

# Docker Compose Commands

---

- Build all containers in YAML file  
`$ docker-compose build`
- Build and run all containers in YAML file:  
`$ docker-compose up`
- List all containers in YAML file:  
`$ docker-compose images`
- Run a one-off command from a container:  
`$ docker-compose run <container-name> <command>`
- Example (after build):  
`$ docker-compose up`

---

# Demo 4: Docker Compose (python)

# Questions?

---

# Additional Resources

---

- Docker: <https://www.docker.com/>
- Docker Docs: <https://docs.docker.com/>
- Docker Hub: <https://hub.docker.com/>

# Thank you!

---

- Please fill out the survey: <http://tinyurl.com/curc-survey18>
- Contact information: [rc-help@Colorado.edu](mailto:rc-help@Colorado.edu)
- Slides:  
[https://github.com/ResearchComputing/Containers\\_Spring\\_2022](https://github.com/ResearchComputing/Containers_Spring_2022)